


StarAi: Deep Reinforcement Learning



Overview

- Recap
- Rainbow Overview
- Double DQN
- Dueling DQN
- Multistep DQN
- Prioritised Experience Replay
- Noisy + Distribution (overview)
- PYSC2

RL environment

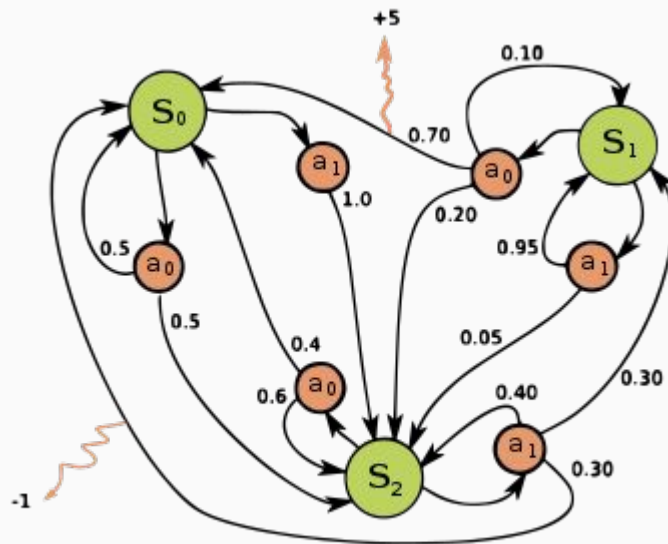


* Image from taken from somewhere on the web :)



Recap

- Epsilon \rightarrow Greedy
- Markov Decision Processes



Recap

- Q Learning (simplified)

$$Q(S_t, A_t) \leftarrow R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$$

- DQN


$$Q(S_t, A_t, \theta) \leftarrow R_{t+1} + \gamma \max_a Q(S_{t+1}, a, \theta_{target})$$

$$L(\theta) = Q(S_t, A_t, \theta) - [R_{t+1} + \gamma \max_a Q(S_{t+1}, a, \theta_{target})]$$

Rainbow



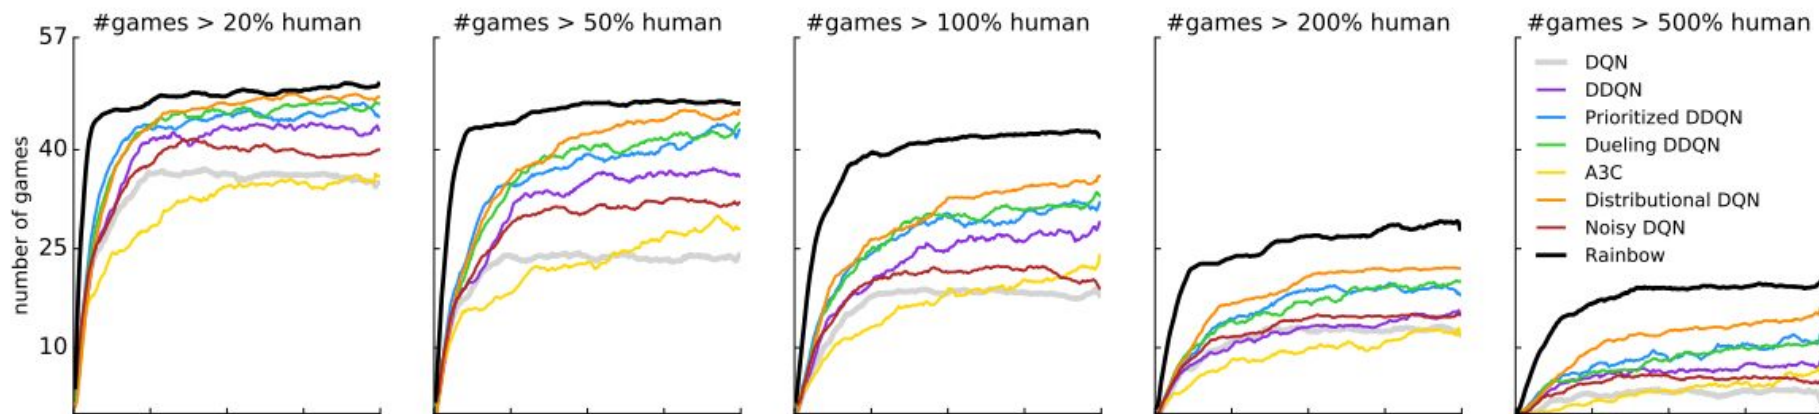
Rainbow on Atari



* Image from taken from deepmind paper




Rainbow on Atari



* Image from taken from deepmind paper




Rainbow Ablations



* Image from taken from deepmind paper



Rainbow Ablations



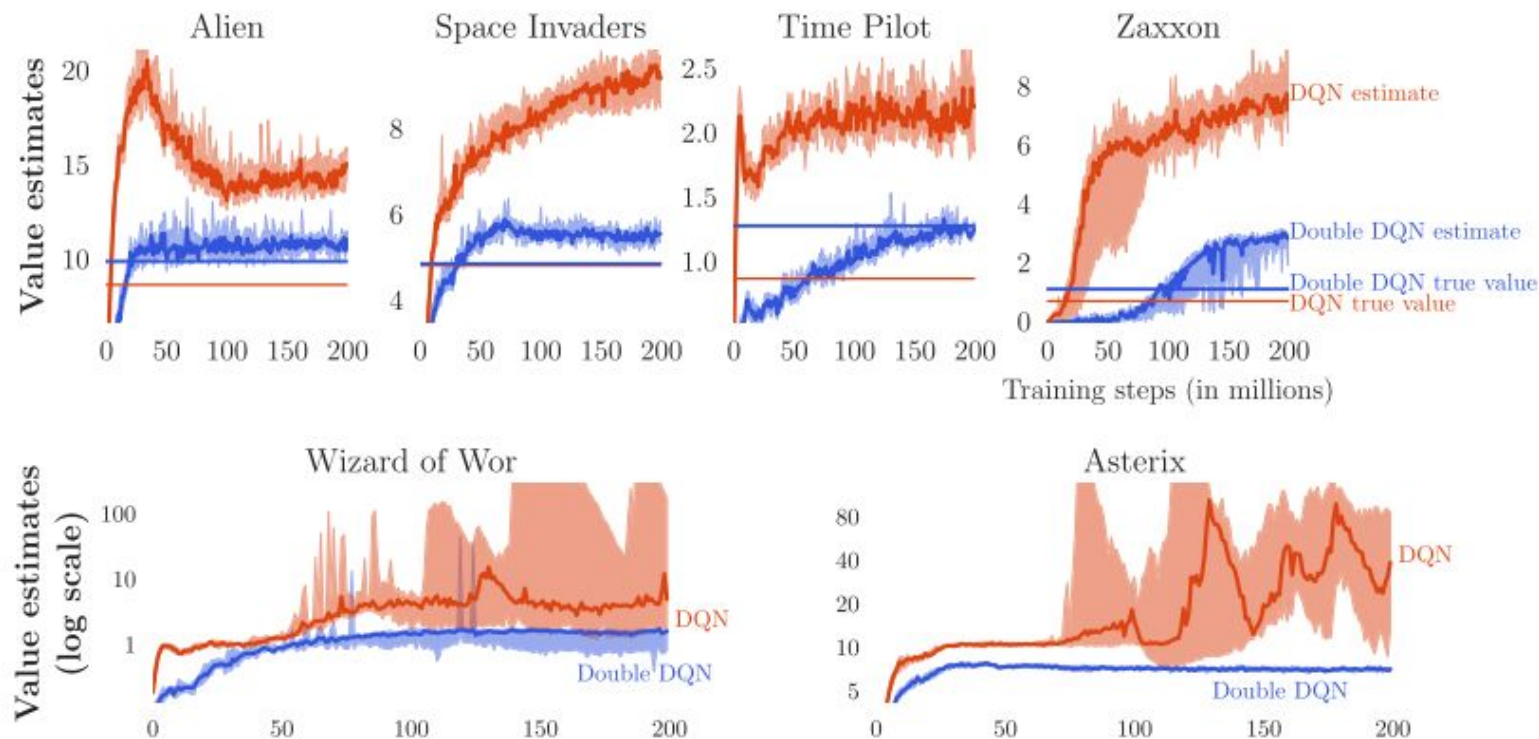
* Image from taken from deepmind paper



Double DQN



Double DQN



* Image from taken from paper



DQN

$$Q(S_t, A_t, \theta) \leftarrow R_{t+1} + \gamma \max_a Q(S_{t+1}, a, \theta_{target})$$

$$L(\theta) = Q(S_t, A_t, \theta) - y \text{ where } y = [R_{t+1} + \gamma \max_a Q(S_{t+1}, a, \theta_{target})]$$

Double Q

everything is the same except for

$$y = [R_{t+1} + \gamma Q(S_{t+1}, \max_a Q(S_{t+1}, a, \theta), \theta_{target})]$$



Link to notebooks


<https://bit.ly/2srCfoH>

Dueling DQN




Dueling DQN

DQN



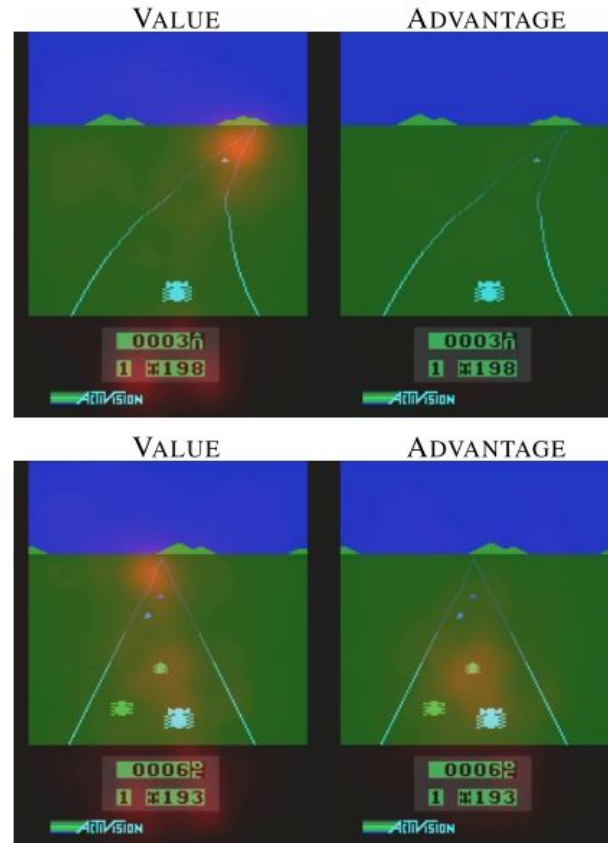
Dueling



* Image from taken from paper



Dueling DQN



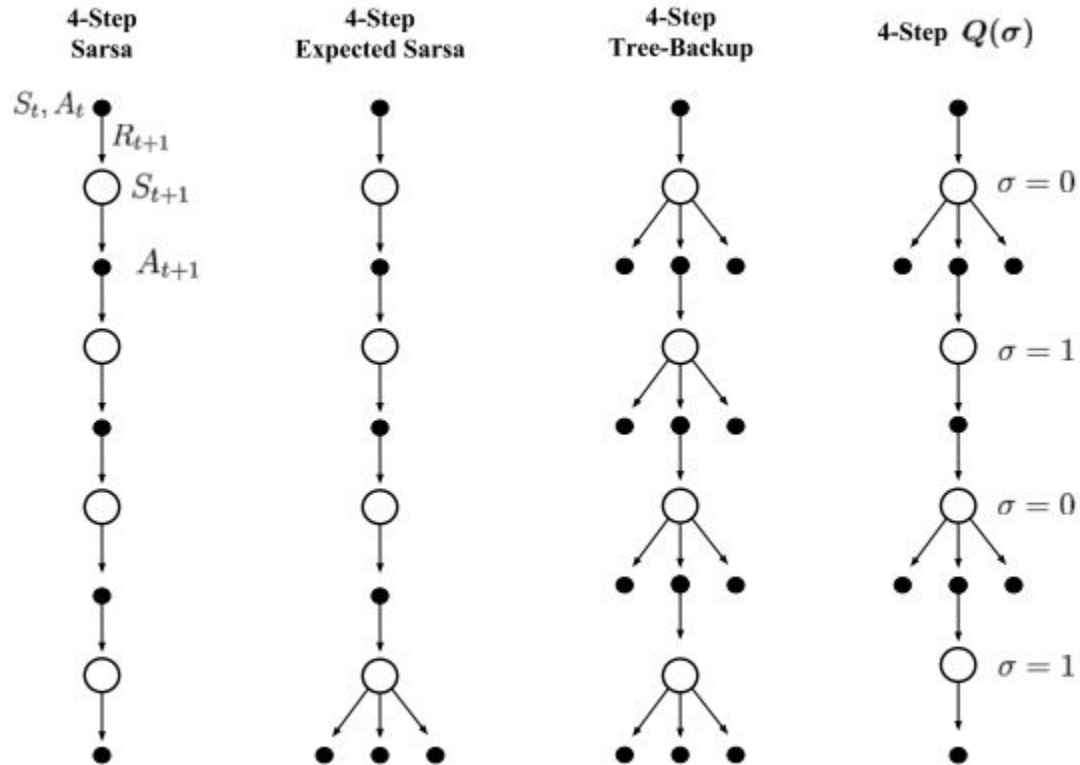
* Image from taken from paper



Multistep



Multistep DQN





* Image from taken from paper



Prioritised Experience Replay



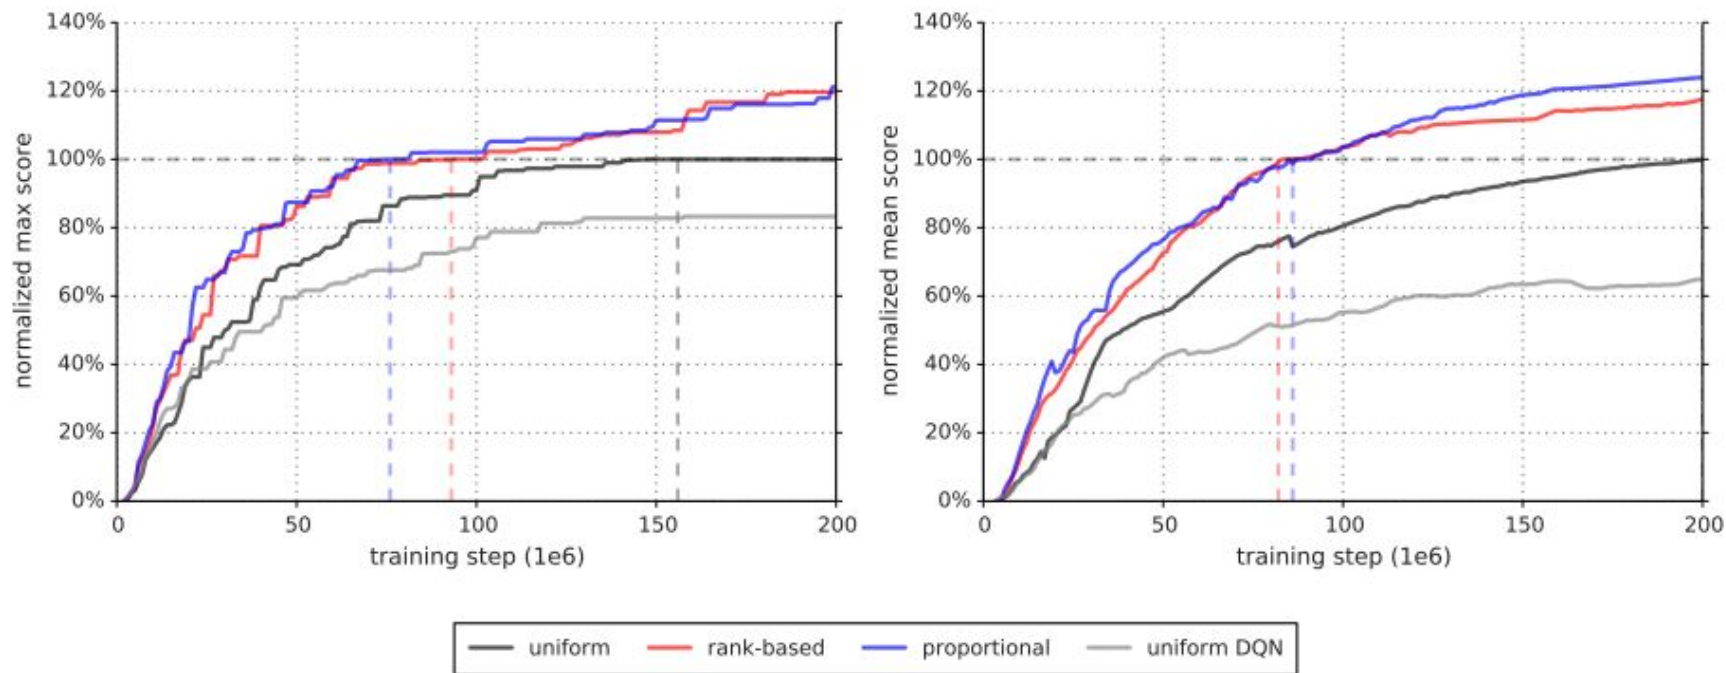
Prioritised Experience Replay - Blind Cliffwalk



* Image from taken from paper



Prioritised Experience Replay - Atari Games

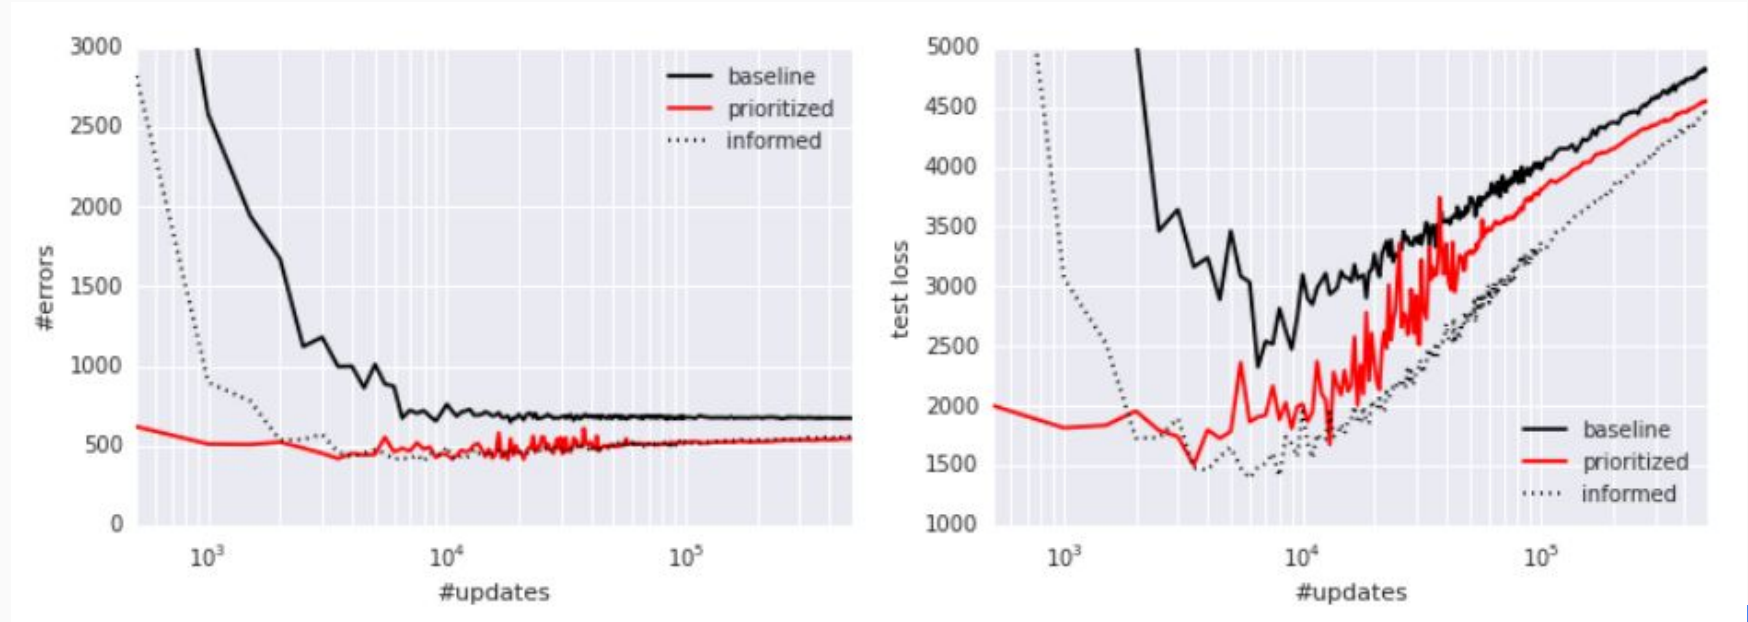


* Image from taken from paper



Prioritised Experience Replay - Supervised Learning

- MNIST
- Removed 99% of samples of 0,1,2,3,4



* Image from taken from paper



Noisy + Distribution



PYSC2



Starcraft 2

- Real time strategy game released in July 2010 by Blizzard and continuously updated
- 3 unique races requiring different tactics to win
- Economy, base management, army management
- Fog of war, exploration, multi tasking and micro management

Watch a match segment



Mini Games

- Move to beacon
- Collect Mineral Shards
- Find and defeat zergling
- Defeat roaches
- Defeat zerglings and banelings
- Collect minerals and gas
- Build marines

Observation Space

Minimap feature layers:

- **height_map**: Shows the terrain levels.
- **visibility**: Which part of the map are hidden, have been seen or are currently visible.
- **creep**: Which parts have zerg creep.
- **camera**: Which part of the map are visible in the screen layers.
- **player_id**: Who owns the units, with absolute ids.
- **player_relative**: Which units are friendly vs hostile. Takes values in $[0, 4]$, denoting [background, self, ally, neutral, enemy] units respectively.
- **selected**: Which units are selected.

Observation Space

Screen feature layers:

- **height_map**: Shows the terrain levels.
- **visibility**: Which part of the map are hidden, have been seen or are currently visible.
- **creep**: Which parts have zerg creep.
- **power**: Which parts have protoss power, only shows your power.
- **player_id**: Who owns the units, with absolute ids.
- **player_relative**: Which units are friendly vs hostile. Takes values in $[0, 4]$, denoting [background, self, ally, neutral, enemy] units respectively.

Observation Space

Screen feature layers:

- **unit_type**: A unit type id, which can be looked up in `pysc2/lib/units.py`.
- **selected**: Which units are selected.
- **hit_points**: How many hit points the unit has.
- **energy**: How much energy the unit has.
- **shields**: How much shields the unit has. Only for protoss units.
- **unit_density**: How many units are in this pixel.
- **Unit_density_aa**: Anti-aliased

Observation Space

Structured layers:

- **General Player Information:** player_id, food, minerals etc
- **Control Groups:** saved unit group info
- **Multi Select:** info about selected unit
- **Cargo:** Units in a transport
- **Build Queue:** Units in production
- **Available actions:** all actions made successfully since last observation
- **Action Result:** result of action
- **Alerts:** whether you are being attacked

Action Space



See it in pycharm

