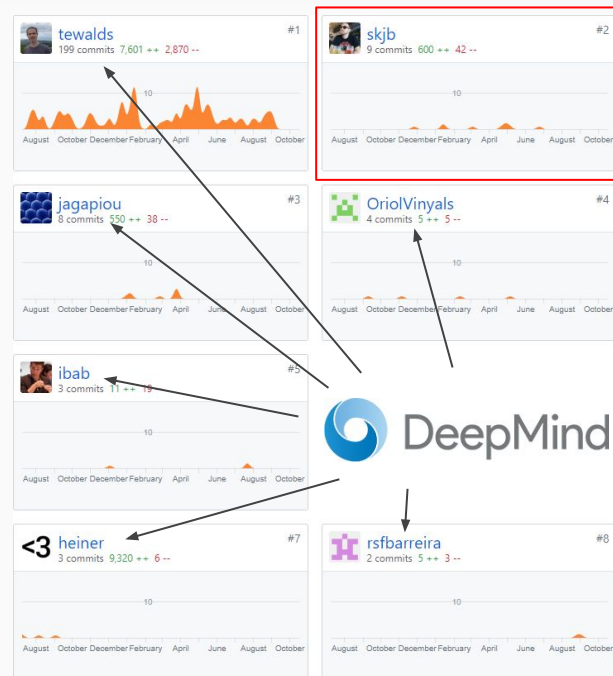


# Introduction to PySC2



# Steven Brown - Who am I?

- 2nd Highest Code Contributor to PySC2
- Creator of Feature Units




# Steven Brown - Where am I?

[medium.com/@skjb](https://medium.com/@skjb)

[twitch.tv/skjb](https://twitch.tv/skjb)

[patreon.com/skjb](https://patreon.com/skjb)



# What is PySC2?

- StarCraft II machine learning environment
- Interacts with Blizzard's StarCraft II API
- Python based
- Designed to emulate human abilities
- Backed by DeepMind
- [github.com/deepmind/pysc2](https://github.com/deepmind/pysc2)

## What is Blizzard's StarCraft II API?

- Allows you to interact with the game via Protobuf
- Has limited Linux support
- Can play replays
- Provides the ability to investigate game state
- Provides the ability to perform player actions
- Supports 2 players
- Has some limitations
- Is still being developed



# Alternatives to PySC2

- Dave Churchill's CommandCenter - C++ framework for BW and SC2, very popular and used for bot battles
- Python SC2 - less human realistic
- C#, Clojure, Java, Go

# Why PySC2?

- Python = TensorFlow and Scikit
- DeepMind development team
- You want to build a human-comparable bot
- You can't stand the thought of interacting with other humans or having any spare time, and you want nothing more than to be coding bots all day

# Why StarCraft II?

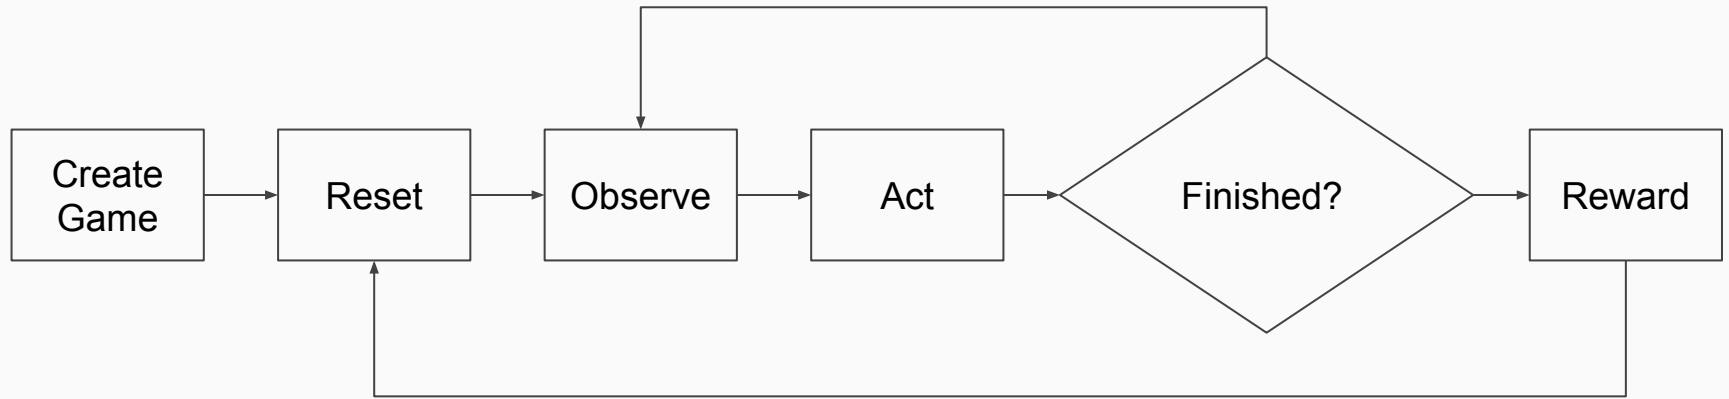
- Deterministic - there's no random chance, in the same state the same action will produce the same result
- Extremely Large State and Action Space - it takes a lot of time to explore and find what works and what doesn't
- Hidden Information - there is information about your opponent you don't know, exacerbated by fog of war



# What sort of agent do you want to build?

- Completely scripted
- Partially scripted, partially ML
- Completely ML but structured specifically for SC2
- Completely ML with no SC2 specific structures

# Game Engine Flow



Remind you of anything?

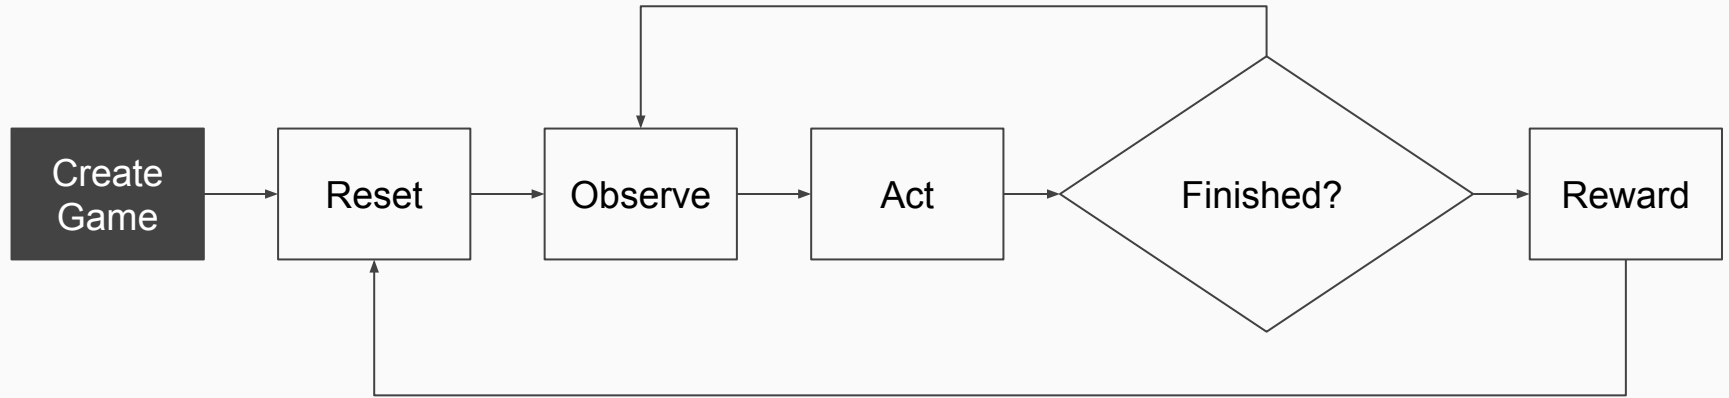


# Example Agent


```
class MyAgent(base_agent.BaseAgent):  
    def __init__(self):  
        super(MyAgent, self).__init__()  
        # One-time setup  
  
    def reset(self):  
        super(MyAgent, self).reset()  
        # Before each game  
  
    def step(self, obs):  
        super(MyAgent, self).step(obs)  
  
        # Read state from obs and ALWAYS act  
        return actions.FUNCTIONS.no_op()
```



# Game Engine Flow



# Main Game Options



## Code Example of Game Options

```
with sc2_env.SC2Env(  
    map_name="Simple64",  
    players=[sc2_env.Agent(sc2_env.Race.terran),  
              sc2_env.Bot(sc2_env.Race.zerg, sc2_env.Difficulty.easy)],  
    agent_interface_format=features.AgentInterfaceFormat(  
        feature_dimensions=features.Dimensions(screen=84,  
                                                minimap=64),  
        action_space=actions.ActionSpace.FEATURES),  
    ) as env:
```



### Simplified games designed for testing algorithms

- BuildMarines
- CollectMineralsAndGas
- CollectMineralShards
- DefeatRoaches
- DefeatZerglingsAndBanelings
- FindAndDefeatZerglings
- MoveToBeacon



- Designed for full game play
- Empty128
- Flat32/48/64/96/128 - No terrain variations
- Simple64/96/128 - Some terrain





- The usual ladder maps
- Can lag behind a bit, depending on Linux support and code updates
- Far more complex to generalise for due to variations in terrain, base locations, resources, etc.



- very\_easy
- easy
- medium
- medium\_hard = Hard
- hard = Harder
- harder = Very hard
- very\_hard = Elite
- cheat\_vision
- cheat\_money
- cheat\_insane



### `step_mul`

The number of steps to take before acting again, there are 22.4 steps per second for normal “faster” games, so a value of 8 = 168 APM, I have had issues below a step mul of 2

### `game_steps_per_episode`

The maximum steps to take before the game automatically ends



### `visualize`

Shows a custom rendered version of the observations, good for debugging but slows the game considerably

### `disable_fog`

Disables the fog of war so everything is visible



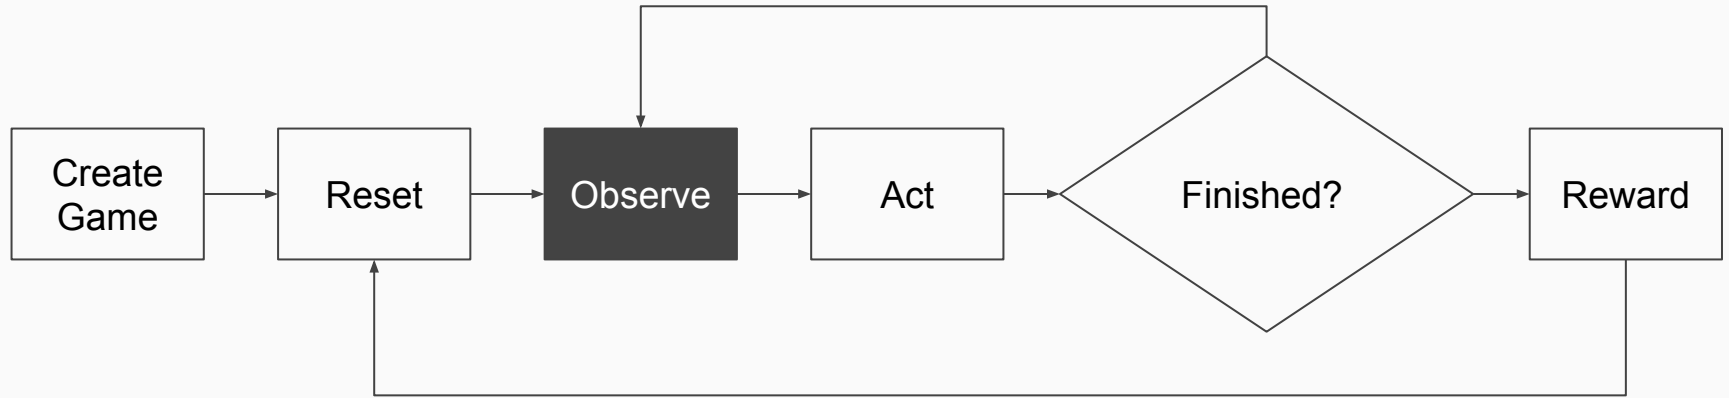
# Machine Learning Tips

- Start against very easy bots
- Play as a single race
- Play against a single race
- Disable the fog of war
- Use a single, simple map
- Human players can beat very easy with an APM <60 so consider a step\_mul of 20-30

# Machine Learning Tips

- Limit the games to half an hour or so (40,320 steps)
- Consider having your agent play against a simple scripted bot
- Consider self-play, you get twice the learning per game but it may not generalise
- Compare your bot against one that chooses completely random actions

# Game Engine Flow



## Example State

```
state = (command_center_count,  
        supply_depot_count,  
        barracks_count,  
        scv_count,  
        marine_count,  
        base1_enemy_count,  
        base2_enemy_count,  
        base3_enemy_count,  
        base4_enemy_count,  
        base1_friendly_count,  
        base2_friendly_count,  
        base3_friendly_count,  
        base4_friendly_count)
```





## Step “obs”

`obs.first()`

Whether or not this is the first step of the game, good for doing things like position and race detection

`obs.last()`

Whether or not this is the last step, use this to learn, save, and wrap things up

`obs.observation.reward`

Use this for sparse rewards, will be 0 for mid-steps or a draw, -1 for a loss, 1 for a win



## Step “obs”

`obs.observation.game_loop`

The current step, you can work out the current second if you divide by 22.4

`obs.observation.available_actions`

Contains a list of actions that can be performed in the current state

`obs.observation`

Everything else



# Main “Features”

Player

Screen

Minimap



Multi-select



Single-select



Build Queue


\*Cargo not shown




# Some Screen Features




Height map



Unit type



Visibility



Player relative


Other important features include power, creep and effects




# Some Minimap Features



Height map



Camera location



Visibility



Player relative

Other important features include creep



# Scalar Feature Layers

x  
↓

y →

[	[	0	0	0	0	0	0	0	0]
[	[	0	0	0	0	0	0	0	0]
[	[	0	255	255	255	212	212	0	0]
[	[	0	255	255	255	212	212	0	0]
[	[	0	0	212	0	212	0	0	0]
[	[	0	212	212	255	255	255	0	0]
[	[	0	212	212	212	255	255	0	0]
[	[	0	0	0	0	0	0	0	0]]



```
height_map = obs.observation.feature_minimap.height_map
height_map[y][x]
height_map[2][1] = 255
```



# Categorical Feature Layers

```
[[0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0]
 [0 3 1 0 3 3 0 0]
 [0 3 1 0 0 0 0 0]
 [0 0 0 0 0 0 0 0]
 [0 3 0 0 0 3 0 0]
 [0 3 3 0 3 3 0 0]
 [0 0 0 0 0 0 0 0]]
```



```
player_relative = obs.observation.feature_minimap.player_relative
```

Remember not to use ranges (e.g. 0-3) for categorical layers when feeding into neural networks, instead supply each category as 0 or 1.



# Categorical Feature Layers

```
[[False False False False False False False False]
 [False False False False False False False False]
 [False False  True False False False False False]
 [False False  True False False False False False]
 [False False False False False False False False]
 [False False False False False False False False]
 [False False False False False False False False]
 [False False False False False False False False]]
```



```
player_relative = obs.observation.feature_minimap.player_relative
minimap_self = (player_relative == features.PlayerRelative.SELF)
```





## Categorical Feature Layers

```
([2, 2], [2, 3]) # [y1, y2], [x1, x2]
```



```
player_relative = obs.observation.feature_minimap.player_relative  
player_y, player_x = (player_relative == features.PlayerRelative.SELF).nonzero()
```



# Categorical Feature Layers

```
([2, 2], # [x1, y1]  
 [3, 2]) # [x2, y2]
```




```
player_relative = obs.observation.feature_minimap.player_relative  
player_y, player_x = (player_relative == features.PlayerRelative.SELF).nonzero()  
player_xy = zip(player_x, player_y)
```



# Screen Feature Notes

- The perspective is different, so don't expect everything to perfectly match the normal game
- Units can overlap and be difficult to identify



## Minimap Feature Notes

- The minimap in the API may not match the minimap in the game
- Each pixel on the minimap can only contain one detail per feature layer



## Enable RGB Observations

```
with sc2_env.SC2Env(  
    map_name="Simple64",  
    players=[sc2_env.Agent(sc2_env.Race.terran),  
              sc2_env.Bot(sc2_env.Race.zerg, sc2_env.Difficulty.easy)],  
    agent_interface_format=features.AgentInterfaceFormat(  
        rgb_dimensions=features.Dimensions(screen=84,  
                                             minimap=64),  
        action_space=actions.ActionSpace.RGB),  
    ) as env:
```



# RGB Observations

```
[[[255 255 255]
  [ 0  0  0]
  [ 0  0  0]
  [ 0  0  0]
  [255 255 255]
  [ 0  0  0]
  [241 191 126]
  [241 191 126]
  [241 191 126]
  [ 0  0  0]
  [ 15 12 11]
  [ 17 11 10]]]
```

B G R



```
obs.observation.rgb_minimap
obs.observation.rgb_screen
```

\*BGR instead of RGB may be a bug



- If you are observing in the RGB space you should act in the RGB space to maintain perspective
- Sizing does not seem to be exact (e.g. specifying 8x8 produced a 12x10 grid)
- Screen perspective matches the regular game
- Minimap seems to match the regular game



# Single, Multi-Select and Cargo Observations

```
obs.observation.single_select  
obs.observation.multi_select  
obs.observation.cargo
```



```
[45  
1  
45  
0  
0  
0  
0]
```

Unit type (SCV)  
Player relative (Self)  
Health  
Shields  
Energy  
Transport slots taken  
Build progress (Not applicable to SCV, normally 0-100)





# Player Observations

```
obs.observation.player.player_id  
obs.observation.player.minerals  
obs.observation.player.vespene  
obs.observation.player.food_used  
obs.observation.player.food_cap  
obs.observation.player.food_army  
obs.observation.player.food_workers  
obs.observation.player.idle_worker_count  
obs.observation.player.army_count  
obs.observation.player.warp_gate_count  
obs.observation.player.larva_count  
  
free_supply = food_cap - food_used
```



## Enable Feature Units

```
with sc2_env.SC2Env(  
    map_name="Simple64",  
    players=[sc2_env.Agent(sc2_env.Race.terran),  
              sc2_env.Bot(sc2_env.Race.zerg, sc2_env.Difficulty.easy)],  
    agent_interface_format=features.AgentInterfaceFormat(  
        feature_dimensions=features.Dimensions(screen=84,  
                                                minimap=64),  
        action_space=actions.ActionSpace.FEATURES,  
        use_feature_units=True),  
    ) as env:
```



- Every visible unit on screen
- Exact unit location
- Build progress
- Assigned worker count
- Ideal worker count

```
marines = [unit for unit in obs.observation.feature_units  
            if unit.unit_type == units.Terran.Marine  
            and unit.alliance == features.PlayerRelative.SELF]
```



# Feature Unit Properties

```
unit_type  
alliance  
health  
shield  
energy  
cargo_space_taken  
build_progress # 0-100  
health_ratio # 0-255  
shield_ratio # 0-255  
energy_ratio # 0-255  
display_type  
owner  
x  
y
```

```
facing  
radius  
cloak  
is_selected  
is_blip  
is_powered  
mineral_contents  
vespene_contents  
cargo_space_max  
assigned_harvesters  
ideal_harvesters  
weapon_cooldown  
order_length  
addon_unit_type # soon?
```



- Unit coordinates may be outside the screen since they are the centre of the unit, you will have to clip the values



- Unit visibility seems to match the real game, so if you are acting in the FEATURE space you may be able to move them to a location that makes them no longer visible



# Machine Learning Tips

- Be sure to use categorical encoding for categorical feature layers, instead of ranges

```
[0 1 0 2 2 0 3 0]  
[0 1 0 0 0 0 0 0]  
[0 0 0 1 1 0 0 0]  
[0 0 0 0 0 0 1 0]
```

```
0 1 2 3
```

# Machine Learning Tips

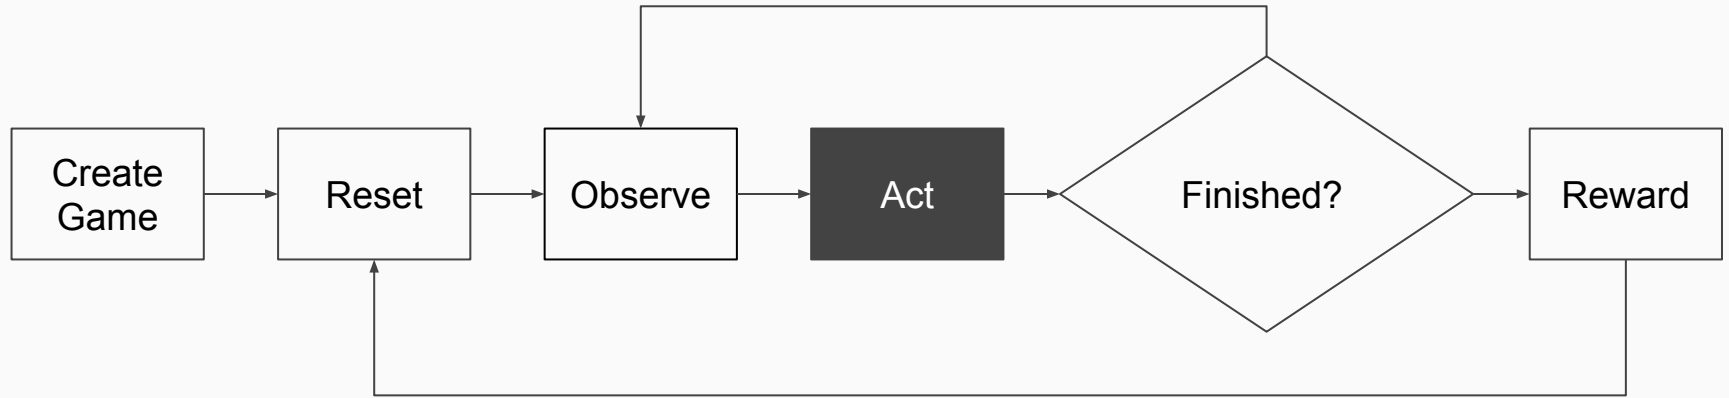
- Feature scaling - things like minerals and step count can number into tens of thousands, while other values may be  $< 10$
- Consider reducing minerals to “can afford” flags
- Crop and rotate the minimap so every game is from the same perspective



# Machine Learning Tips

- Make sure the agent can tell the difference between distinct states - don't leave out crucial information

# Game Engine Flow



# Move the Camera


Get the camera position:

```
camera_ys, camera_xs = (obs.observation.feature_minimap.camera == 1).nonzero()
camera_y = camera_ys.mean()
camera_x = camera_xs.mean()
camera_size = camera_y.max - camera_y.min # assuming square resolution
```

Set the camera position:

```
return actions.FUNCTIONS.move_camera((x, y)) # (x, y) is a tuple
```

- Due to rounding errors (or something) the camera position you receive may not match what you send in
- The x and y coordinates must be within the minimap size, e.g. 0-63
- Top-left is (0, 0)



## Select an Idle Worker

```
if obs.observation.player.idle_worker_count > 0:  
    return actions.FUNCTIONS.select_idle_worker("select")
```

Selecting idle workers will move the screen

You can also “select\_all”



## Build a Barracks

```
if (actions.FUNCTIONS.Build_Barracks_screen.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Build_Barracks_screen("now", (x, y))
```

You can help your agent by preventing it from building in the mineral line

For Protoss or Zerg agents you can help by limiting to power or creep



## Build a Barracks

```
if (actions.FUNCTIONS.Build_Barracks_screen.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Build_Barracks_screen("now", (x, y))
```

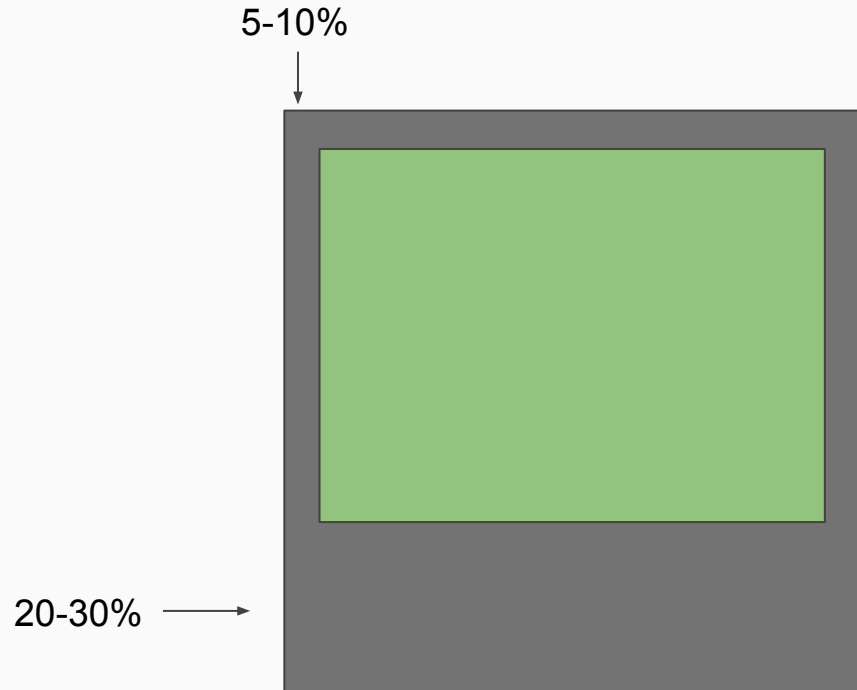
You can help your agent by preventing it from building in the mineral line

For Protoss or Zerg agents you can help by limiting to power or creep



# Building Locations

It's best to add a margin to the left, top, and right, and a larger margin to the bottom



```
if (actions.FUNCTIONS.Effect_Repair_autocast.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Effect_Repair_autocast()
```





# Harvest Minerals

```
if (actions.FUNCTIONS.Harvest_Gather_screen.id in
    obs.observation.available_actions):
    minerals = [unit for unit in self.obs.observation.feature_units
                 if unit.unit_type in MINERAL_UNIT_TYPES]
    if len(minerals) > 0:
        mineral = random.choice(minerals)

    return actions.FUNCTIONS.Harvest_Gather_screen(
        "queued", (mineral.x, mineral.y))
```

I have not had a lot of success with queued mineral harvesting



# Mineral Unit Types

```
units.Neutral.BattleStationMineralField  
units.Neutral.BattleStationMineralField750  
units.Neutral.LabMineralField  
units.Neutral.LabMineralField750  
units.Neutral.MineralField  
units.Neutral.MineralField750  
units.Neutral.PurifierMineralField  
units.Neutral.PurifierMineralField750  
units.Neutral.PurifierRichMineralField  
units.Neutral.PurifierRichMineralField750  
units.Neutral.RichMineralField  
units.Neutral.RichMineralField750
```



## Vespene Unit Types

```
units.Neutral.ProtossVespeneGeyser  
units.Neutral.PurifierVespeneGeyser  
units.Neutral.RichVespeneGeyser  
units.Neutral.ShakurasVespeneGeyser  
units.Neutral.SpacePlatformGeyser  
units.Neutral.VespeneGeyser
```



## Selecting all Barracks

```
barracks = [unit for unit in self.obs.observation.feature_units
             if unit.unit_type == units.Terran.Barracks
             and unit.alliance == features.PlayerRelative.SELF]
if len(barracks) > 0:
    barrack = random.choice(barracks)

    return actions.FUNCTIONS.select_point(
        "select_all_type", (barrack.x, barrack.y))
```

Remember to clip the x and y coordinates to fit the screen resolution (e.g. 0-83)

Barrack is not the singular for barracks, I know



## Set Rally Point

```
if (actions.FUNCTIONS.Rally_Units_screen.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Rally_Units_screen("now", (x, y))
```

You can also rally workers

You can also use the minimap



## Add Barracks to Control Group

```
return actions.FUNCTIONS.select_control_group("append", 0)
```

You can also set and recall

I have had issues with this



# Train a Marine

```
if (actions.FUNCTIONS.Train_Marine_quick.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Train_Marine_quick("now")
```

There seems to be a bug currently that stops the correct distribution of build orders

In my experience when you have multiple units selected, all commands apply to all units even if normally they would not (e.g. SCVs and buildings)



## Upgrade to Orbital Command

```
if (actions.FUNCTIONS.Morph_OrbitalCommand_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Morph_OrbitalCommand_quick("now")
```





```
if (actions.FUNCTIONS.Effect_Scan_screen.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Effect_Scan_screen("now", (x, y))
```

You can also scan the minimap



```
if (actions.FUNCTIONS.Build_TechLab_quick.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Build_TechLab_quick("now")
```

There is currently a bug that only allows “Build\_Techlab\_screen”, so you need to lift with “Lift\_Barracks\_quick” and then use “Build\_Techlab\_screen” with coordinates



- By default common actions are grouped, e.g. “burrow” for Zerg units has an individual command for each unit, but when they are grouped you only need to issue one command
- You can disable this by setting `hide_specific_actions=False` in your agent interface format



```
if (actions.FUNCTIONS.Research_Stimpack_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Research_Stimpack_quick("now")
```

It's hard to know if your research is done since research build queues don't exist (yet?), you may have to internally track how long it has been in progress, and check again if the research action is available (and you have enough resources)



```
if (actions.FUNCTIONS.Effect_Stim_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Effect_Stim_quick("now")
```



# Attack

```
if (actions.FUNCTIONS.Attack_minimap.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Attack_minimap("now", (x, y))
```

You can also attack the screen

You can also move units without attacking



# Machine Learning Tips

- Consider scripted action sequences
- Consider scripted building locations
- Consider scripted build orders
- Attacking is simpler if you choose the closest enemy location
- Consider automated supply management
- Consider automated worker management

# Machine Learning Tips

- Try limiting actions only to those that are valid in the current state
- If limiting actions, remember to exclude those actions when learning

```
[-0.1, 0.0] # max is 0.0, should be -0.1  
[ 0.0, 0.0] # max is 0.0
```

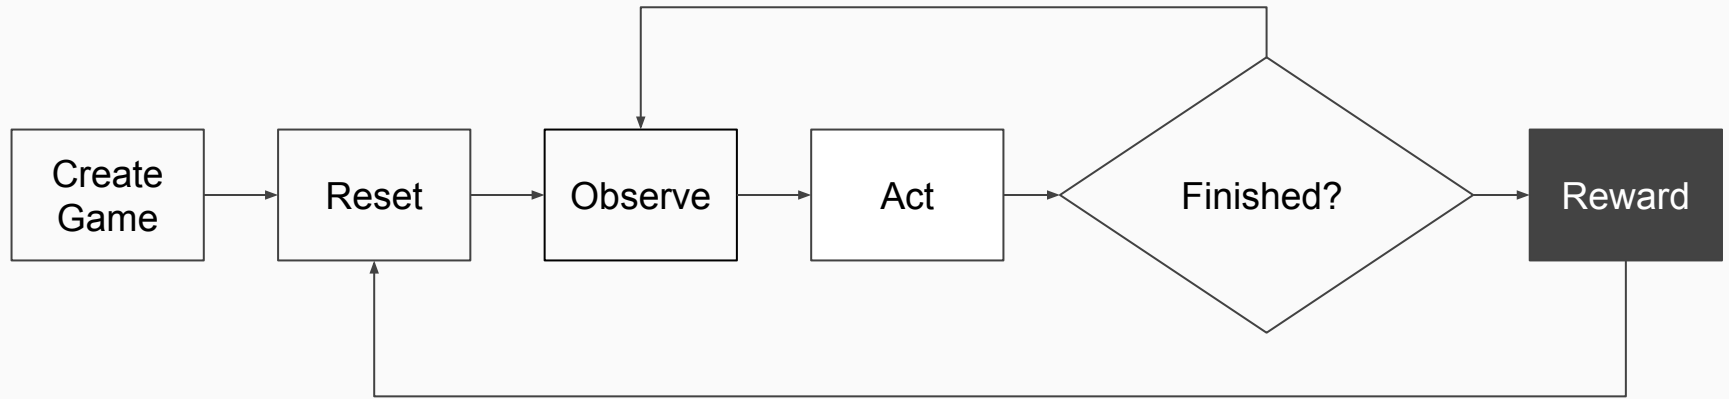


# Machine Learning Tips

- Do not learn if the state does not change (unless it's the last step) as actions will tend towards the highest value even if they had no impact

```
[ 0.0, -0.1, 0.1]
[ 0.01, -0.1, 0.1] # action 0
[ 0.0, -0.008, 0.1] # action 1
```

# Game Engine Flow



```
if obs.last():  
    reward = obs.reward  
  
    # train your agent  
    # save your learning  
  
return actions.FUNCTIONS.no_op()
```



# Machine Learning Tips

- You can define your own reward structure
- Be careful what you reward - you just might get it
- You might split your agent and reward for different things, for example:
  - Building units
  - Killing units
  - Finding the enemy

## Limitations

- Cannot accept surrender
- Hard to track upgrades
- No research progress
- No ghosted buildings
- No build queue in multi-select
- Alerts/messages seem to be limited and unreliable



- Can you play against your bot?
- Does it support supervised learning?
- Can it process replays?



- Linux is used by DeepMind, so that is most thoroughly tested
- Linux releases lag behind the latest balance changes and ladder maps
- Linux bugs take longer to get fixed
- Windows stays updated as it's the official game
- PySC2 often lags behind the latest Windows balance changes and ladder maps



# More Resources

- <https://discord.gg/b2gjyHR>
- <https://medium.com/@skjb/build-a-zerg-bot-with-pysc2-2-0-295375d2f58e>